Uncovering Bugs in Distributed Storage Systems during Testing (not in Production!)

Pantazis Deligiannis[†], Matt McCutchen[◇], Paul Thomson[†], Shuo Chen[⋆]
Alastair F. Donaldson[†], John Erickson[⋆], Cheng Huang[⋆], Akash Lal[⋆]
Rashmi Mudduluru[⋆], Shaz Qadeer[⋆], Wolfram Schulte[⋆]

[†]*Imperial College London,* [◇]*Massachusetts Institute of Technology,* [⋆]*Microsoft*

Abstract

Testing distributed storage systems is challenging due to multiple sources of nondeterminism. Conventional testing techniques, such as unit and stress testing, are ineffective in preventing serious but subtle bugs from reaching production. Formal techniques, such as TLA+, can only verify high-level specifications of systems at the level of logic-based models, and fall short of checking the actual executable code. In this paper, we present a new methodology for testing distributed storage systems. Our approach applies advanced systematic testing techniques to thoroughly check that the executable code adheres to its high-level specifications, which significantly improves coverage of important system behaviors.

Our methodology has been applied to three distributed storage systems in the Microsoft Azure cloud computing platform. In the process, numerous bugs were identified, reproduced, confirmed and fixed. These bugs required a subtle combination of concurrency and failures, making them extremely difficult to find with conventional testing techniques. An important advantage of our approach is that a bug is uncovered in a small setting and witnessed by a full system trace, which dramatically increases the productivity of debugging.

1 Introduction

Distributed systems are notoriously hard to design, implement and test [5, 24, 20, 26]. This challenge is due to many well-known sources of *nondeterminism* [6, 18], such as unexpected node failures, the asynchronous interaction between system components, data losses due to unreliable communication channels, the use of multithreaded code to exploit multicore machines, and interaction with clients. All these sources of nondeterminism can easily create *Heisenbugs* [13, 28], corner-case bugs that are difficult to detect, diagnose and fix. A Heisenbug might hide deep inside one of these paths and only manifest under extremely rare conditions [13, 28], but the consequences can be catastrophic [1, 33].

Developers of production distributed systems use many testing techniques, such as unit testing, integration testing, stress testing, and fault injection. In spite of extensive use of these testing methods, many bugs that arise from subtle combinations of concurrency and failure events are missed during testing and get exposed only in production. However, allowing serious bugs to reach production can cost organizations a lot of money [31] and lead to customer dissatisfaction [1, 33].

We interviewed technical leaders and senior managers in Microsoft Azure regarding the top problems in distributed system development. The consensus was that one of the most critical problems today is how to improve *testing coverage* so that bugs can be uncovered *during testing* and *not in production*. The need for better testing techniques is not specific to Microsoft; other companies, such as Amazon and Google, have acknowledged [6, 29] that testing methodologies have to improve to be able to reason about the correctness of increasingly more complex distributed systems that are used in production.

Recently, the Amazon Web Services (AWS) team used formal methods "to prevent serious but subtle bugs from reaching production" [29]. The gist of their approach is to extract the high-level logic from a production system, represent this logic as specifications in the expressive TLA+ [22] language, and finally verify the specifications using a model checker. While highly effective, as demonstrated by its use in AWS, this approach falls short of "verifying that executable code correctly implements the high-level specification", and the AWS team admits that it is "not aware of any such tools that can handle distributed systems as large and complex as those being built at Amazon" [29].

We have found that checking high-level specifications is necessary but not sufficient, due to the gap between the specification and the executable code. Our goal is to bridge this gap. We propose a new methodology that

validates high-level specifications directly on the executable code. Our methodology is different from prior approaches that required developers to either switch to an unfamiliar domain specific language [19, 8], or manually annotate and instrument their code [34]. Instead, we allow developers to test *production code* by writing test harnesses in C#, a mainstream programming language. This significantly lowered the acceptance barrier for adoption by the Microsoft Azure team.

Our methodology is based on P# [7], an extension of C# that provides support for modeling, specification, and systematic testing of distributed systems written in Microsoft's .NET framework. To test a distributed system with P#, the programmer augments the original system code with three artifacts – a model of the nondeterministic execution environment of the system, a concurrent test harness that drives the system towards interesting behaviors, and safety or liveness specifications. The P# testing engine then systematically explores all behaviors exercised by the test harness and validates them against the provided specifications.

We have applied P# to three distributed storage systems inside Microsoft: Azure Storage vNext; Azure Table Live Migration; and Azure Service Fabric. We uncovered numerous bugs in these systems, including, most notably, a subtle liveness bug that only intermittently manifested during stress testing for months without being fixed. Our testing approach uncovered this bug in a very small setting, which made it easy for developers to examine traces, identify, and fix the problem.

To summarize, our contributions are as follows:

- We present a new methodology that allows flexible modeling of the environment of a distributed system using simple language mechanisms in P#.

- Because P# is an extension of C#, our infrastructure can test production code written in C#.

- We present three case studies of using P# to test production distributed systems, finding bugs that could not be found with traditional testing techniques.

- We show that P# can reproduce bugs in production-scale distributed systems in a small setting with easy to understand traces.

2 Motivating Example

Microsoft Azure Storage is a cloud storage system that provides customers the ability to store seemingly limitless amounts of data. It has grown from tens of petabytes in 2010 to exabytes in 2015, with the total number of objects stored well-exceeding 60 trillion [14].


Figure 1: Top-level components for extent management in Microsoft Azure Storage vNext.

Azure Storage vNext is the next generation storage system for Microsoft Azure, where the primary design target is to scale the storage capacity by more than $100\times$. Similar to the current system, vNext employs containers, called *extents*, to store data. Extents are typically several gigabytes each, consisting of many data blocks, and replicated over multiple *Extent Nodes* (ENs). However, in contrast to the current system, which employs a Paxos-based, centralized mapping from extents to ENs, vNext achieves its scalability target by employing a *distributed mapping*. In vNext, extents are divided into partitions, with each partition managed by a light-weight *Extent Manager* (ExtMgr). This partitioning is illustrated in Figure 1.

One of the many responsibilities of an ExtMgr is to ensure that every extent maintains enough *replicas* in the system. To achieve this, an ExtMgr receives frequent periodic *heartbeat* messages from every EN. The failure of an EN is detected by missing heartbeats. An ExtMgr also receives less frequent, but still periodic, *synchronization reports* from every EN. The sync reports list all the extents (and associated metadata) stored on the EN. Based on these two types of messages, an ExtMgr identifies which ENs have failed, and which extents are affected by the failure and are missing replicas as a result. The ExtMgr then schedules tasks to repair the affected extents and distributes the tasks to the necessary ENs. The ENs then repair the extents from their existing replicas in the system and lazily update the ExtMgr via their next periodic sync reports. All this communication between an ExtMgr and the ENs occurs via network engines installed in each component of vNext (see Figure 1).

To ensure correctness, the developers of vNext have instrumented extensive, multiple levels of testing:

1. *Unit testing*, in which emulated heartbeats and sync reports are sent to an ExtMgr. These tests check that the messages are processed as expected.

2. *Integration testing*, in which an ExtMgr is launched

2

together with multiple ENs. An EN failure is subsequently injected. These tests check that the affected extents are eventually repaired.

3. *Stress testing*, in which an ExtMgr is launched together with multiple ENs and multiple extents. The test keeps repeating the following process: injecting an EN failure, launching a new EN and checking that the affected extents are eventually repaired.

Despite of the extensive testing efforts, the vNext developers have been plagued for months by an elusive bug in the ExtMgr logic. All the unit test suites and integration test suites successfully pass every single time. However, the stress test suite fails *from time to time* after very long executions; in these cases, certain replicas of some extents fail and are never repaired. This bug has proven to be difficult to identify, reproduce and troubleshoot. First, it requires very long executions to trigger. Second, an extent not being repaired is *not* a property that can be easily verified. In practice, the developers rely on a very large time-out period to detect the bug. Finally, by the time that the bug is detected, very long execution traces have been collected, which makes manual inspection tedious and ineffective.

To uncover this bug, as well as many other tricky bugs, the vNext developers are in constant search of a generic and systematic approach for testing production-scale distributed storage systems.

3 Our Approach to Testing Distributed Systems

The goal of our work is to develop testing techniques for detecting and debugging Heisenbugs in distributed storage systems prior to deployment. To achieve our goal, we use P# [7], an extension of the mainstream C# language that provides: (i) language support for *modeling the environment* of distributed systems developed using Microsoft's .NET framework; and (ii) a *testing engine* that can systematically explore all interleavings between asynchronous event handlers, as well as other nondeterministic events such as failures and timeouts.

A P# program consists of multiple *state machines* that communicate with each other *asynchronously* by sending and receiving *events*, which may contain an optional *payload*. A P# machine declaration is similar to a class declaration in C#: it can contain an arbitrary number of fields and methods. However, a machine declaration differs from a class declaration, as it also contains an input event queue, and one or more *states*. Each state can register *actions* to handle incoming events.

P# machines run concurrently with each other, each executing an event handling loop that dequeues an event from the input queue and handles it by invoking the registered action. This action might access a field, call a method, transition the machine to a new state, create a new machine, or send an event to another machine. In P#, a send operation is *non-blocking*; the event is simply enqueued into the input queue of the target machine, which will dequeue and handle the event concurrently. All this functionality is provided in a lightweight runtime library, built on top of Microsoft's Task Parallel Library [25].

Our approach of using P# to test distributed systems, implemented in .NET, requires the developer to perform the following three key modeling tasks. These tasks are illustrated in §4 for the Azure Storage vNext system.

1. The real components of a distributed system must be wrapped inside P# machines. In addition, the environment of the system must be modeled using P#, so all external nondeterminism can be captured. We collectively refer to the environment model of a system and its wrapped components as the P# *test harness*. This test harness can be used by P# to drive the system towards interesting behaviors.

2. The criteria for correctness of an execution must be specified. Specifications in P# can encode either *safety* or *liveness* [21] properties (see §3.1). Safety specifications generalize the notion of source code assertions; a safety violation is a finite trace leading to an erroneous state. Liveness specifications generalize nontermination; a liveness violation is an infinite trace that exhibits lack of progress.

3. All the asynchrony due to message passing in the system must be modeled using the P# event-sending APIs. Likewise, all other sources of nondeterminism, such as failures and timer expiration, must be modeled using appropriate P# APIs. This step allows P# to systematically explore all asynchronous event handler interleavings and all other sources of nondeterminism during testing (see §3.2).

P# uses modern object-oriented language features such as *interfaces* and *virtual method dispatch* to connect the real code with the modeled code. Programmers in industry are used to working with such features, and heavily use them in production for testing purposes. In our experience, this significantly lowers the bar for product teams inside Microsoft to embrace P# for testing.

In principle, our modeling methodology is *not specific* to P# and the .NET framework, and can be used in combination with any other programming framework that has equivalent capabilities. We also argue that our approach is *flexible* since it allows the user to model *as much* or *as little* of the environment as required to achieve the desired level of testing.

3.1 Specifications

In this section, we describe how safety and liveness specifications are expressed in P#.

Safety specifications. In addition to the usual assertions for stating safety properties that are local to a machine, P# also provides a way to specify global assertions by using a *safety monitor* [8], which is a special machine that can receive, but not send, events.

A safety monitor maintains local state that is modified in response to events received from ordinary (non-monitor) machines. This local state is used to maintain a history of the computation that is relevant to the property being specified. An erroneous global behavior is flagged via an assertion on the private state of the safety monitor. Thus, a P# monitor cleanly separates instrumentation state required for specification (inside the monitor) from program state (outside the monitor).

Liveness specifications. Liveness property specifications are more difficult to encode since they are intended to capture program progress over infinite executions. Violation of a liveness property can only be demonstrated by an infinite execution. Usually, a liveness property is specified via a temporal logic formula [30, 22]. We take a different approach and allow the programmer to write a *liveness monitor* [8]. Similar to a safety monitor, a liveness monitor can receive, but not send, events. Unlike a safety monitor, a liveness monitor contains two types of states: *hot* and *cold*.

A hot state denotes a point in the execution where progress is required but has not happened yet; for example, a node has failed but a new one has not come up yet. A liveness monitor enters a hot state upon receiving notification of an event that requires the system to make progress. The liveness monitor leaves the hot state and enters a cold state upon receiving another event notifying it of progress. An infinite execution is erroneous if the liveness monitor is in a hot state infinitely often but in a cold state only finitely often. Our liveness monitors can encode arbitrary temporal logic properties. A thermometer analogy is useful for understanding liveness monitors: a hot state increases the temperature by a small value; a cold state resets the temperature to zero; a liveness error happens if the temperature becomes infinite.

A liveness violation is witnessed by an *infinite* execution in which all concurrently executing P# machines are *fairly* scheduled. Unfortunately, it is clearly not possible to generate an infinite execution by executing a program for a finite amount of time. Therefore, our implementation of liveness checking in P# approximates an infinite execution using several heuristics. One such heuristic considers an execution longer than a large user-supplied bound as an "infinite" execution [19, 27]. Another heuristic maintains a cache of (pieces of) the state of the P# program obtained at each step in the execution, and reports an "infinite" execution when the latest step results in a cache hit, thus approximating a cycle in the state graph of the program.

3.2 Systematic testing

The P# runtime is a lightweight layer build on top of the Task Parallel Library (TPL) of .NET that implements the semantics of P#: creating state machines, executing them concurrently using the default task scheduler of TPL, and sending events and enqueueing them in the appropriate machines. A key capability of the P# runtime is that it can execute in *bug-finding* mode, which systematically tests a P# program to find bugs using the safety and liveness monitors that were discussed in §3.1.

When the P# runtime executes in bug-finding mode, an embedded *systematic concurrency testing engine* [12, 28, 10] captures and takes control of all sources of nondeterminism that are *known* to the P# runtime. In particular, the runtime is aware of nondeterminism due to the interleaving of event handlers in different machines. Each send operation to an ordinary (non-monitor) machine, and each create operation of an ordinary machine, creates an interleaving point where the runtime makes a scheduling decision. The runtime is also aware of all nondeterministic events (e.g. failures) that were captured during the modeling process.

The P# testing engine will repeatedly execute a program from start to completion, each time exploring a potentially different set of scheduling decisions, until it either reaches a bound (in number of iterations or time), or it hits a property violation. This testing process is fully automatic, has no false-positives (assuming an accurate environment model), and can reproduce found bugs by replaying buggy schedules. After a bug is found, the engine generates a trace that represents the buggy schedule. In contrast to logs typically generated during production, the P# trace provides a global order of all communication events and, thus, is easier to debug.

We have implemented two different schedulers inside the P# systematic testing engine: a *random* scheduler that makes each nondeterministic choice randomly, and a *probabilistic concurrency testing* [4] (PCT) scheduler. It is straightforward to create a new scheduler by implementing the ISchedulingStrategy interface [9] exposed by the P# libraries. The interface exposes callbacks that are invoked by the P# testing engine for taking decisions regarding which machine to schedule next, and can be used for developing both generic and application-specific schedulers.

Handling C# 5.0 asynchrony. Our case studies heavily use the async and await C# 5.0 language primitives. An async method is able to perform await on a TPL

4

task *without* blocking the current thread. The C# compiler achieves this non-blocking behavior by wrapping the code following an `await` statement in a TPL *task continuation*, and then returning execution to the caller of the `async` method. The TPL scheduler executes this task continuation when the task being awaited has completed. In principle, this asynchrony can be modeled using the P# primitives for machine creation and message passing. However, in our experience, modeling async/await asynchronous code using P# is difficult and time consuming, and thus infeasible for large code bases. To tackle this problem, we developed a solution that automatically captures asynchronously created TPL tasks and wraps them in temporary P# machines.

Our solution involves using a *custom task scheduler*, during systematic testing with P#, instead of the default TPL scheduler. All TPL tasks spawned in the P# program will be scheduled for execution in the P# task scheduler. These tasks will be intercepted by our custom task scheduler and wrapped inside a special (not available to the programmer) task-machine. When the P# testing engine schedules a task-machine for execution, the corresponding task is unwrapped and executed. This approach works nicely with the `async` and `await` C# 5.0 primitives: when an `async` method is called, the asynchronous TPL task will be automatically wrapped in a task-machine and scheduled by P#; for `await`, no special treatment is required since the task continuation will also be automatically wrapped and scheduled accordingly.

4 Testing Azure Storage vNext with P#

To uncover the elusive extent repair bug in Azure Storage vNext, its developers wrote a testing harness in P#. The developers of vNext expected that it was more likely for the bug to occur in the ExtMgr logic, rather than the EN logic. Hence, the testing harness focuses on the real ExtMgr using mocked ENs.

The testing harness consists of the following P# machines (as shown in Figure 2):

ExtentManager acts as a thin wrapper machine for the real ExtMgr component in vNext.

ExtentNode represents a simplified mocked EN.

TestingDriver communicates with all other machines, relays messages between the `ExtentManager` and the `ExtentNode` machines, and is responsible for driving testing scenarios.

RepairMonitor collects states from all `ExtentNode` machines in the testing harness and checks desired properties, such as that an extent is replicated or repaired at the expected ENs.
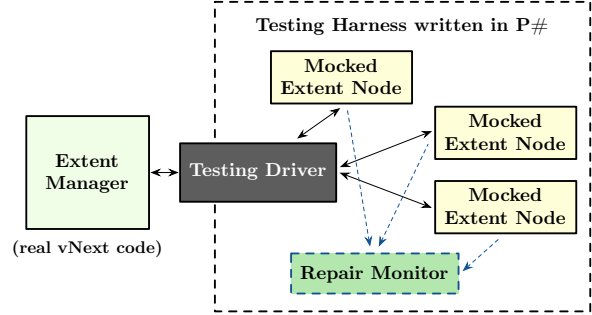

Figure 2: Real Extent Manager with a mocked environment (each box represents one P# machine).

4.1 The ExtentManager machine

The real ExtMgr in vNext, which is our system-under-test, is wrapped inside the `ExtentManager` P# machine in the testing harness.

Internals of the real Extent Manager. Inside the real ExtMgr (see Figure 3), there are two data structures related to extent replication and repair: `ExtentCenter` and `ExtentNodeMap`. The `ExtentCenter` data structure maintains the mapping records from extents to their hosting ENs. It is updated upon the periodic sync reports from the ENs. Recall that the sync report from a particular EN lists all the extents stored at the EN. Its purpose is to update extent manager's possible out-of-date view of the EN with the ground truth. `ExtentNodeMap` records the latest heartbeat time from every EN.

The ExtMgr runs internally a periodic *EN expiration loop* that is responsible for removing ENs that have been missing heartbeats for an extended period of time, as well as cleaning up the corresponding extent records in `ExtentCenter`. In addition, the ExtMgr runs a periodic *extent repair loop* that examines all the `ExtentCenter` records, identifies extents with missing replicas, schedules extent repair tasks and sends them to the ENs.

Intercepting messages between machines. The real ExtMgr uses a network engine to send messages to the ENs. The testing harness mocks the original network engine in vNext and overrides its interface. In this way, the mocked network engine intercepts all outbound messages and relays them to the `TestingDriver` machine, which is responsible for dispatching the messages to the corresponding `ExtentNode` machines. As shown in Figure 4, the mocked network engine intercepts the outbound messages from the ExtMgr and invokes `PSharpRuntime.Send(...)` to asynchronously relay the messages to `TestingDriver`. Conceivably, the mocked network engine could leverage the nondeterminism support in P# and choose to drop the messages in a non-deterministic fashion, in case emulating message loss is desirable.
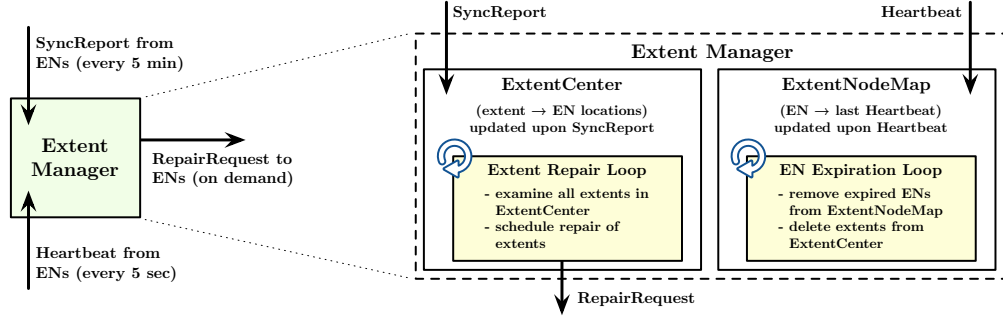
Figure 3: Internal components of the real Extent Manager in Azure Storage vNext.

```
// network interface in vNext
class NetworkEngine {
  public virtual void SendMessage(Socket s, Message msg);
}

// mocked engine for intercepting Extent Manager messages
class MockedNetEngine : NetworkEngine {
  public override void SendMessage(Socket s, Message msg) {
    // intercept and relay Extent Manager messages
    PSharpRuntime.Send(this.TestingDriver,
      new MessageFromExtentManagerEvent(), s, msg);
  }
}
```

Figure 4: Mocked network engine in vNext.

Specifics of the ExtentManager machine. Figure 5 shows code snippets from the `ExtentManager` P# machine. `ExtentManager` is a thin wrapper of the real ExtMgr. The mocked network engine replaces the real one in ExtMgr, intercepts all the outbound messages from ExtMgr and relays them to `TestingDriver`.

Messages coming from the ENs do *not* go through the mocked network engine. They are delivered to the `ExtentManager` machine directly and trigger an action that invokes the messages on the internal ExtMgr with `extMgr.ProcessMessage`. The benefit of this approach is that ExtMgr can be tested without modifying its code; ExtMgr is simply unaware of the testing harness and behaves as if it running in a real distributed environment and communicating with real ENs.

System correctness should *not* hinge on the frequency of any individual timer. Instead, all nondeterminism due to timing-related events must be delegated to P#. To achieve this, all timers inside ExtMgr are disabled, and the EN expiration loop and the extent repair loop are driven instead by timers modeled in P#, an approach that was used in previous work [8]. The P# timers send nondeterministic timeout events that are controlled by the P# runtime. Hence, the P# testing engine has the freedom to schedule arbitrary interleavings between timeout events and all other regular system events.

```
// Wrapping the target vNext component in a P# machine
class ExtentManagerMachine : Machine {
  private ExtentManager extMgr; // real vNext code

  void Init() {
    extMgr = new ExtentManager();
    extMgr.netEngine = new MockedNetEngine(); // mock network
    extMgr.isMockingTimer = true; // disable internal timer
  }

  [OnEvent(MessageFromExtentNode, DeliverExtentNodeMessage)]
  void DeliverExtentNodeMessage() {
    var msg = (ExtentNodeMessage)this.Payload;
    // relay messages from Extent Node to Extent Manager
    extMgr.ProcessMessage(msg);
  }

  [OnEvent(TimerTick), nameof(ProcessExtentRepair))]
  void ProcessExtentRepair() {
    // extent repair loop driven by external timer
    extMgr.ProcessEvent(new ExtentRepairEvent());
  }
}
```

Figure 5: System-under-test: instance of the real Extent Manager wrapped inside the `ExtentManager` machine.

4.2 The ExtentNode machine

The `ExtentNode` machine is a simplified version of the original EN. The machine omits much of the complex details of a real EN, and only mocks the logic necessary for the testing scenarios. This mocked logic includes: repairing an extent from its replica, and sending sync reports and heartbeat messages periodically to the `ExtentManager` machine.

The testing harness leverages components of Azure Storage vNext whenever it is appropriate. For example, the `ExtentNode` machine re-uses the `ExtentCenter` data structure, a component that is used inside a real EN for extent bookkeeping.

In the mocked extent repair logic, `ExtentNode` takes action upon receiving an extent repair request from the `ExtentManager` machine. It sends a copy request to a source `ExtentNode` machine where a replica is stored. After receiving an `ExtentCopyResponse` event from the source, it updates the internal `ExtentCenter`, as il-

```
// Mocking Extent Node in P#
class ExtentNodeMachine : Machine {
  // leverage real vNext component whenever appropriate
  private ExtentNode.ExtentCenter extCtr;

  // extent repair logic
  ...
  [OnEvent(ExtentCopyResponse, ProcessExtentCopyResponse)]
  void ProcessExtentCopyResponse() {
    // extent copy response from source replica
    if (IsCopySucceeded(this.Payload)) {
      var rec = GetExtentRecord(this.Payload);
      extCtr.AddOrUpdate(rec); // update ExtentCenter
    }
  }

  // extent node sync logic
  [OnEvent(TimerTick, ProcessExtentNodeSync)]
  void ProcessExtentNodeSync() {
    var sync = extCtr.GetSyncReport(); // prepare sync report
    PSharpRuntime.Send(this.ExtentManagerMachine,
      new MessageFromExtentNodeEvent(), sync);
  }
}
```

Figure 6: The mocked EN in vNext.

```
class RepairMonitor : Monitor {
  // true: EN has replica, false: EN has no replica
  private Dictionary<Machine, bool> ExtentNode2Replica;

  // cold state: repaired
  [OnEvent(NotifyEnFailure, ProcessExtentNodeFailure)]
  cold state Repaired {
    void ProcessExtentNodeFailure() {
      var node = GetExtentNode(this.Payload);
      ExtentNode2Replica.Remove(node);
      goto Repairing;
    }
  }

  // hot state: repairing
  [OnEvent(NotifyExtentRepaired, ProcessExtentRepaired)]
  hot state Repairing {
    void ProcessExtentRepaired() {
      var node = GetExtentNode(this.Payload);
      ExtentNode2Replica[node] = true;
      if (ReplicaCount == Harness.REPLICA_COUNT_TARGET)
        goto Repaired;
    }
  }
}
```

Figure 7: The RepairMonitor machine.

lustrated in Figure 6.

In the mocked EN sync logic, the machine is again driven by an external timer provided by P#. It prepares a sync report with `extCtr.GetSyncReport(...)` and asynchronously sends the report to `ExtentManager` with `PSharpRuntime.Send(...)`.

4.3   The TestingDriver machine

The `TestingDriver` machine drives two testing scenarios. In the first scenario, it launches one `ExtentManager` and three `ExtentNode` machines, with a single extent on one of the ENs. It waits for the extent to be replicated at the other ENs. In the second scenario, it fails one of the `ExtentNode` machines and launches a new one. It waits for the extent to be repaired on the new EN. We omit illustrative code snippets due to space constraints.

4.4   The RepairMonitor machine

The `RepairMonitor` machine is a P# liveness monitor (see §3.1) that transitions between a cold and a hot state. Whenever an EN fails, `RepairMonitor` is notified. As soon as the number of extent replicas falls below a specified target (three replicas in the current P# test harness), `RepairMonitor` transitions into the hot *repairing* state, where the missing replica is being repaired. Whenever a replica is repaired, `RepairMonitor` is notified. It transitions into the cold *repaired* state, when the replica number reaches again the target, as illustrated in Figure 7.

In the extent repair testing scenarios, `RepairMonitor` checks that it should *always eventually* end up in the cold state. Otherwise, `RepairMonitor` is stuck in the hot state for *infinitely* long. This indicates that the corre-

sponding execution sequence results in an extent replica never being repaired, which is a liveness bug.

4.5   Liveness Bug in Azure Storage vNext

It took less than ten seconds for the P# testing engine to report the first occurrence of a liveness bug in vNext (see §6). Upon examining the debug trace, the developers of vNext were able to quickly confirm the bug.

However, the original P# trace did not include sufficient detail to allow the developers to identify the root cause of the problem. Fortunately, running the test harness took very little time, so the developers were able to quickly iterate and add more refined debug outputs in each iteration. After several iterations, the developers were able to pinpoint the exact culprit and immediately propose a solution for fixing the bug. Once the proposed solution was implemented, the developers ran again the test harness. No bugs were found during 100,000 iterations, a process that required only tens of minutes.

The liveness bug occurs in the second testing scenario, where the `TestingDriver` machine fails one of the `ExtentNode` machines and launches a new one. `RepairMonitor` transitions to the hot repairing state and is stuck in the state for infinitely long.

The following is one particular execution sequence resulting in this liveness bug: (i) EN_0 fails and is detected by the EN expiration loop; (ii) EN_0 is removed from `ExtentNodeMap`; (iii) `ExtentCenter` is updated and the replica count drops from 3 (which is the target) to 2; (iv) ExtMgr receives a sync report from EN_0; (v) the extent center is updated and the replica count increases again from 2 to 3. This is problematic since the replica count

is equal to the target, which means that the extent repair loop will never schedule any repair task. At the same time, there are only two *true* replicas in the system, which is one less than the target. This execution sequence leads to one missing replica; repeating this process two more times would result in all replicas missing, but ExtMgr still thinking that all replicas are healthy. In production, such a bug can cause a very serious incident of customer data loss.

The culprit is in (iv), where ExtMgr receives a sync report from EN_0 after deleting the EN. This may occur in P# due to arbitrary interleaving of events. It may also occur, albeit much less frequently, during stress testing due to messages being delayed in the network. This explains why the bug only occurs from time to time during stress testing and requires long executions to manifest. In contrast, P# allows the bug to manifest quickly, the developers to iterate rapidly, the culprit to be identified promptly, and the fix to be tested effectively and thoroughly, all of which have the potential to vastly increase the productivity of distributed storage system development.

5 Additional Case Studies

The modeling and testing approach described in the earlier sections of the paper is not specific to the vNext system. P# is a generic framework for .NET distributed systems. We showcase this capability by presenting developer experience of using P# to model and test two other systems used in production inside Microsoft.

5.1 Live Azure Table Migration

The Live Azure Table Migration library (MigratingTable) is designed to transparently migrate a data set between tables in an Azure Storage service while an application is accessing this data set. This case study is particularly interesting because its developers *co-developed* the system with its P# test to increase confidence in the system. Indeed, this co-development process resulted in the discovery of several tricky bugs before the system went into production (see §6). The MigratingTable testing effort differs from the vNext case study: the vNext developers focused on specifying a single liveness property, while the MigratingTable test checks complete compliance with an interface specification.

MigratingTable provides a *virtual table* that has a similar interface to an ordinary Azure table. This interface is named `IChainTable`. The virtual table is backed by a pair of *old* and *new* tables, which must also implement `IChainTable`. A background *migrator* job is responsible for moving all data from the old table to the new table. Meanwhile, each read and write operation issued to the virtual table is translated to a sequence of reads
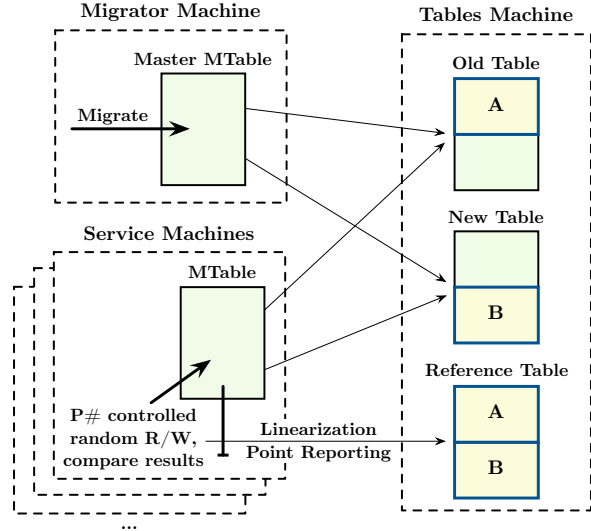

Figure 8: Environment model of MigratingTable (each box with a dotted line represents one P# machine).

and writes on the backend tables according to a protocol, which is intended to guarantee that when multiple application processes issue *input* operations to their own MigratingTable instances with the same backend tables, the output of the MigratingTable instances is compliant with the specification of `IChainTable` for the combined *input history* as if all the operations were done on a single table. The goal of the P# test was to verify this property.

Modeling and testing. The main challenge behind testing MigratingTable is that there are many possible input histories and that the system is highly concurrent. The developers could have chosen specific input histories to test, but they were not confident that this approach would be effective in catching bugs, especially because concurrency increases the potential for difficult-to-foresee interactions between seemingly independent parts of the code. Instead they took advantage of the simplicity and near-determinism of the `IChainTable` specification to write an in-memory reference implementation, called `SpecTable`, to which the output of MigratingTable on an arbitrary input history can be compared. Then the test could sample from a distribution that was defined over all possible input histories within certain bounds, so that under the *small scope hypothesis* that any given bug in MigratingTable leads to incorrect output for at least one input history in the distribution, each iteration of the test would have a positive probability of detecting the bug.

The MigratingTable was instrumented to report the intended *linearization point* of each input operation on the abstract, shared virtual table. This allows the test harness to keep a *reference* `SpecTable` in sync with the

virtual table by sending each operation to the reference table when the linearization point on the virtual table is reached. The output of the operation on the MigratingTable is then checked against its output on the reference table.

The P# test harness of MigratingTable consists of a `Tables` machine containing the old and new backend tables and the reference table; a collection of `Service` machines containing identically configured MigratingTables; and a `Migrator` machine that performs the background migration (see Figure 8). Each `Service` machine issues a random sequence of input operations to its MigratingTable, which sends backend operations to the `Tables` machine. The `Service` machine sends the operation to the reference table at the linearization point and compares the results as previously described.

P# captures and controls the interleaving of the backend operations. To ensure that the reference table is never observed to be out of sync with the backend tables, after the `Tables` machine processes a backend operation, it enters a state that defers further backend operations until MigratingTable has reported whether the backend operation was a linearization point and (if so) the operation on the reference table has been completed.

5.2   Azure Service Fabric

Azure Service Fabric[1] (or Fabric for short) is a platform and API for creating reliable services that execute on a cluster of machines. In order to make the user-written service *reliable*, Fabric launches several *replicas* (copies) of the service, where each replica runs as a separate process on a different node in the cluster. One replica is selected to be the *primary* which serves client requests; the rest are *secondaries*. The primary replicates state changes to the secondaries so that all replicas eventually have the same state. If the primary fails, Fabric elects one of the secondaries to be the new primary and launches another secondary; the new secondary will receive a copy of the state of the new primary in order to "catch up" with the other secondaries. User-written Fabric services are complex asynchronous and distributed applications, and are thus challenging to test.

Our primary goal was to create a P# model of Fabric to allow thorough testing of services, where Fabric's asynchrony is controlled by the P# runtime. The model is written once to include all behaviours of Fabric, including simulating failures and recovery, so that it can be reused repeatedly to test multiple Fabric services. This was the largest modeling exercise among the case studies but the cost was amortized across testing multiple servies. This is analogous to the work on driver verification [2] where the cost of developing a model of the

---
[1] `azure.microsoft.com/campaigns/service-fabric/`

Windows kernel was amortized across testing multiple device drivers. Note that we model the lowest Fabric API layer (`Fabric.dll`) which is currently not documented for use externally. We target internally-developed services that invoke the lowest layer.

P#'s systematic testing was very helpful in debugging the model itself. In order to perform systematic testing on our model, we wrote a simple service in P# that runs on our Fabric model. We tested a scenario where the primary replica fails at some non-deterministic point within the execution. One such bug that we found involved the primary replica failing when a secondary replica was about to receive a copy of the state; the secondary was then elected to be the new primary and yet, because the secondary stopped waiting for a copy of the state, it was also promoted to be an *active* secondary (one that has caught up with the other secondaries). This caused an assertion failure in our model allowing us to detect and fix this incorrect behaviour.

The main system that we tested is *CScale* [11], a big data-stream processing system. In prior work [7], we used an earlier version of the Fabric model and reported bugs in sample services. Supporting CScale required significant additions to the model, making it much more feature-complete. CScale chains multiple Fabric services, which communicate via remote procedure calls (RPCs). To close the system, we modeled RPCs by sending and receiving P# events. This was enough to converted a distributed system that uses both Fabric and its own network communication protocol into a closed single process system. A key challenge in our work was to test CScale despite the fact that it uses various synchronous and asynchronous APIs other than P#. This work is still in-progress. However, we were able to find a `NullReferenceException` bug in CScale by running it on our model. The bug has been communicated to the developers of CScale, but we are still awaiting a confirmation.

6   Quantifying the Cost of Using P#

We report our experience of applying P# on the three case studies discussed in this paper. We aim to answer the following two questions:

1. How much human effort was spent in modeling the environment of a distributed system using P#?

2. How much computational time was spent in systematically testing a distributed system using P#?

6.1   Cost of environment modeling

Environment modeling is a core activity of using P#. It is required for *closing* a system to make it amenable to

System-under-test	System		P# Test Harness			
	#LoC	#B	#LoC	#M	#ST	#AH
vNext Extent Manager	19,775	1	684	5	11	17
MigratingTable	2,267	11	2,275	3	5	10
Fabric User Service	31,959	1*	6,534	13	21	87

Table 1: Statistics from modeling the environment of the three Microsoft Azure-based systems under test. The ($\star$) denotes "awaiting confirmation".

systematic testing. Table 1 presents program statistics for the three case studies. The columns under "System" refer to the real system-under-test, while the columns under "P# Test Harness" refer to the environment model written in P#. We report: lines of code for the system-under-test (#LoC); number of bugs found in the system-under-test (#B); lines of P# code for the test harness (#LoC); number of machines (#M); number of state transitions (#ST); and number of action handlers (#AH).

Modeling the environment of the Extent Manager in the Azure Storage vNext system required approximately two person weeks of part-time developing. The P# test harness for this system is the smallest (in lines of code) from the three case studies. This was because this modeling exercise aimed to reproduce the particular liveness bug that was haunting the developers of vNext.

Developing both the MigratingTable system and its P# test harness took approximately five person weeks. The P# test harness was developed in parallel with the actual system. This is in contrast with the other two case studies discussed in this paper, where the modeling activity occurred independently and after the development process had finished.

Modeling Fabric required approximately five person months, an effort undertaken by the authors of P#. In contrast the other two systems discussed in this paper were modeled and tested by their corresponding developers. Although modeling Fabric required a significant amount of time, it is a one-time effort, which only needs incremental refinement with each release.

6.2 Cost of systematic testing

Using P# we managed to uncover 8 serious bugs in our case studies. As discussed earlier in the paper, these bugs were hard to find with traditional testing techniques, but P# managed to uncover them and reproduce them in a small setting. According to the developers, the traces of P# were useful, as it allowed them to understand the source of the bug and fix it in a timely manner. After the developers fixed all the discovered bugs, we optionally reintroduced them one-by-one to evaluate the effectiveness of different P# systematic testing strategies in find-

ing these bugs.

Table 2 presents the results from running the P# systematic testing engine on each case study with a re-introduced bug using the random and PCT schedulers. We chose to use controlled random scheduling, because it has proven to be efficient for finding concurrency bugs [32, 7]. The CS column shows which case study corresponds to each bug: "1" is for the Azure Storage vNext; and "2" is for MigratingTable. We do not include the Fabric case study, as we are awaiting confirmation of the found bug (see §5.2).

We performed all experiments using the Windows PowerShell tool on a 2.50GHz Intel Core i5-4300U CPU with 8GB RAM running Windows 10 Pro 64-bit. We configured the engine to perform 100,000 iterations. The random seed for both schedulers was generated using the current time. The PCT scheduler was further configured with a bug depth of 2 and a max number of scheduling steps of 500. All reported times are in seconds.

For the vNext case study, the random scheduler was able to reproduce the bug in less than 10 seconds. The reason that the number of scheduling steps to find the bug is much higher than the rest of the bugs in the table is that this bug is a liveness violation: as discussed in §3.2 we leave the program to run for a long time before checking if the liveness property holds. The PCT scheduler was unable to find the bug using the bug depth of 2, which suggests that the bug requires a larger depth bound to be found.

For the MigratingTable case study, we evaluated the P# test harness of §5.1 on 11 "bugs", including 8 *organic* bugs that actually occurred in development and 3 *notional* bugs (marked $*$), which are other code changes that we deemed interesting ways of making the code incorrect. The test found seven of the organic bugs, which are listed first in Table 2. The remaining four bugs (denoted with $\diamond$) were not caught with our default test harness in the 100,000 iterations (a process that required less than 30 minutes). We believe this is because the inputs and schedules that trigger them are too rare in our distribution. To confirm this, we wrote a custom test case for each bug with a specific input that triggers it and were able to quickly reproduce the bugs, as shown. This method is useful for regression testing, but one could not use it during development with no prior knowledge about the bugs.

The QueryStreamedBackUpNewStream bug in MigratingTable, that was found using P#, stands out because it reflects the type of oversight that tends to occur as designs evolve. We do not discuss the details of the bug due to space constraints, but P# managed to discover this bug in a matter of seconds. The MigratingTable developers spent just 10 minutes analyzing the trace to diagnose what was happening; although, this was

CS	Bug Identifier	P# Random Scheduler			P# PCT Scheduler		
		Time to Bug (s)	#SS	BF?	Time to Bug (s)	#SS	BF?
1	ExtentNodeLivenessViolation	9.83	9,000	✓	-	-	✗
2	QueryAtomicFilterShadowing	157.22	165	✓	350.46	108	✓
2	QueryStreamedLock	2,121.45	181	✓	6.58	220	✓
2	QueryStreamedBackUpNewStream	-	-	✗	5.95	232	✓
2	DeleteNoLeaveTombstonesEtag	-	-	✗	4.69	272	✓
2	DeletePrimaryKey	2.72	168	✓	2.37	171	✓
2	EnsurePartitionSwitchedFromPopulated	25.17	85	✓	1.57	136	✓
2	TombstoneOutputETag	8.25	305	✓	3.40	242	✓
◇2	QueryStreamedFilterShadowing	0.55	79	✓	0.41	79	✓
*◇2	MigrateSkipPreferOld	-	-	✗	1.13	115	✓
*◇2	MigrateSkipUseNewWithTombstones	-	-	✗	1.16	120	✓
*◇2	InsertBehindMigrator	0.32	47	✓	0.31	47	✓

Table 2: Results from running the P# random and PCT systematic testing schedulers for 100,000 iterations. We report: time in seconds to find a bug (Time to Bug); number of scheduling steps when a bug was found (#SS); and if a bug was found with a particular scheduler (BF?).

after days of experience analyzing traces. The developers extended the logging of P# with additional trace information to understand the bug. P# only outputs trace information related to its communicating state machines, but the trace information is easy to extend and was done in all our case studies.

7 Related Work

Model checking [12] is a powerful technique that has been widely used in the past for finding tricky bugs in the actual implementation of distributed systems [19, 37, 36, 15, 17, 34, 24].

State-of-the-art model checkers, such as MODIST [37] and dBug [34], typically focus on testing entire, often unmodified, distributed systems, an approach that easily leads to state-space explosion. In contrast, we try to offer a more pragmatic approach for handling state-space explosion: P# provides the capability for creating models of individual components of a large system, which helps towards component isolation and enables compositional testing. Our approach aims to enhance unit and integration testing, techniques widely used in production, where only individual (or a small number of) components are tested at each time.

Tools such as CrystalBall [36] and DEMETER [17] employ techniques based on dynamic partial order reduction [23] to reduce the exploration state-space. The SAMC [24] model checker offers a way of incorporating application-specific information during systematic testing, but the focus is still on reducing the set of possible interleavings. Although SAMC is orders of magnitude faster than previous testing techniques, the authors mention that their "extensive evaluation exercised more than 100,000 executions and used approximately 48 full machine days" [24]. P# is able to explore a similar number of iterations in a matter of hours, and find serious bugs in seconds (see §6). However, we do *not* claim that P# is faster; as aforementioned, we focus on compositional testing, whereas prior work typically focuses on testing an entire system. We are just trying to solve a similar problem at a very different level.

MACEMC [19] is a model checker for distributed systems implemented in the MACE language. The focus of MACEMC is to find liveness property violations using an algorithm based on bounded random walk and the use of heuristics. P# differs from MACEMC in that it can be applied on legacy code written in a mainstream language, whereas a system to be tested with MACEMC has to be written in MACE, which makes MACEMC harder to be applied in an industrial setting.

FATE and DESTINI is a framework for systematically injecting (combination of) failures in distributed systems [16]. This framework focuses in exercising failure scenarios, whereas P# can be used for testing generic safety and liveness properties of distributed systems. P# also provides language extensions and modeling capabilities that further differentiate it from prior work.

Formal methods have been successfully used in industry to verify the correctness of distributed protocols. A recent notable example is the use of TLA+ [22] by the Amazon Web Services team [29]. TLA+ is an expressive formal specification language that can be used to design and verify concurrent programs via model checking. A limitation of TLA+, as well as other similar specification languages, is that they are applied on a model of the sys-

tem and not the actual system. Even if the model is verified, the gap between the real-world implementation and the verified model is still significant, so implementation bugs are still a realistic concern.

Another formal approach is Verdi [35], a framework for writing and verifying distributed systems in Coq [3]. Verdi generates OCaml code from the verifying system, which can be used for execution. In contrast, P# performs bounded testing on a system already written in C#, which in our experience lowers the bar for adoption by product groups. Furthermore, Verdi does not currently support detecting liveness property violations, an important class of bugs in distributed storage systems.

8 Conclusion

We presented a new methodology for testing distributed storage systems. Our approach involves using P#, an extension of the C# language that provides advanced modeling, specification and systematic testing capabilities. We reported experience on applying our methodology on three distributed storage systems that are used inside Microsoft. Using P#, the developers of these systems found, reproduced, confirmed and fixed numerous bugs.

References

[1] AMAZON. Summary of the AWS service event in the US East Region. http://aws.amazon.com/message/67457/, 2012.

[2] BALL, T., LEVIN, V., AND RAJAMANI, S. K. A decade of software model checking with SLAM. *Communications of the ACM 54*, 7 (2011), 68–76.

[3] BARRAS, B., BOUTIN, S., CORNES, C., COURANT, J., FILLIATRE, J.-C., GIMENEZ, E., HERBELIN, H., HUET, G., MUNOZ, C., MURTHY, C., ET AL. The Coq proof assistant reference manual: Version 6.1. https://hal.inria.fr/inria-00069968/, 1997.

[4] BURCKHARDT, S., KOTHARI, P., MUSUVATHI, M., AND NAGARAKATTE, S. A randomized scheduler with probabilistic guarantees of finding bugs. In *Proceedings of the 15th International Conference on Architectural Support for Programming Languages and Operating Systems* (2010), ACM, pp. 167–178.

[5] CAVAGE, M. There's just no getting around it: you're building a distributed system. *ACM Queue 11*, 4 (2013), 30–41.

[6] CHANDRA, T. D., GRIESEMER, R., AND REDSTONE, J. Paxos made live: an engineering perspective. In *Proceedings of the 26th Annual ACM Symposium on Principles of Distributed Computing* (2007), ACM, pp. 398–407.

[7] DELIGIANNIS, P., DONALDSON, A. F., KETEMA, J., LAL, A., AND THOMSON, P. Asynchronous programming, analysis and testing with state machines. In *Proceedings of the 36th ACM SIGPLAN Conference on Programming Language Design and Implementation* (2015), ACM, pp. 154–164.

[8] DESAI, A., JACKSON, E., PHANISHAYEE, A., QADEER, S., AND SESHIA, S. A. Building reliable distributed systems with P. Tech. Rep. UCB/EECS-2015-198, EECS Department, University of California, Berkeley, Sep 2015.

[9] DESAI, A., QADEER, S., AND SESHIA, S. A. Systematic testing of asynchronous reactive systems. In *Proceedings of the 2015 10th Joint Meeting on Foundations of Software Engineering* (2015), pp. 73–83.

[10] EMMI, M., QADEER, S., AND RAKAMARIĆ, Z. Delay-bounded scheduling. In *Proceedings of the 38th Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages* (2011), ACM, pp. 411–422.

[11] FALEIRO, J., RAJAMANI, S., RAJAN, K., RAMALINGAM, G., AND VASWANI, K. CScale: A programming model for scalable and reliable distributed applications. In *Proceedings of the 17th Monterey Conference on Large-Scale Complex IT Systems: Development, Operation and Management* (2012), Springer-Verlag, pp. 148–156.

[12] GODEFROID, P. Model checking for programming languages using VeriSoft. In *Proceedings of the 24th ACM Symposium on Principles of Programming Languages* (1997), ACM, pp. 174–186.

[13] GRAY, J. Why do computers stop and what can be done about it? In *Proceedings of the 5th Symposium on Reliability in Distributed Software and Database Systems* (1986), IEEE, pp. 3–12.

[14] GREENBERG, A. SDN for the Cloud. Keynote in the 2015 ACM Conference on Special Interest Group on Data Communication, 2015.

[15] GUERRAOUI, R., AND YABANDEH, M. Model checking a networked system without the network. In *Proceedings of the 8th USENIX Conference on Networked Systems Design and Implementation* (2011), USENIX, pp. 225–238.

[16] GUNAWI, H. S., DO, T., JOSHI, P., ALVARO, P., HELLERSTEIN, J. M., ARPACI-DUSSEAU, A. C., ARPACI-DUSSEAU, R. H., SEN, K., AND BORTHAKUR, D. FATE and DESTINI: A framework for cloud recovery testing. In *Proceedings of the 8th USENIX Conference on Networked Systems Design and Implementation* (2011), USENIX, pp. 238–252.

[17] GUO, H., WU, M., ZHOU, L., HU, G., YANG, J., AND ZHANG, L. Practical software model checking via dynamic interface reduction. In *Proceedings of the 23rd ACM Symposium on Operating Systems Principles* (2011), ACM, pp. 265–278.

[18] HENRY, A. Cloud storage FUD: Failure and uncertainty and durability. Keynote in the 7th USENIX Conference on File and Storage Technologies, 2009.

[19] KILLIAN, C., ANDERSON, J. W., JHALA, R., AND VAHDAT, A. Life, death, and the critical transition: Finding liveness bugs in systems code. In *Proceedings of the 4th USENIX Conference on Networked Systems Design and Implementation* (2007), USENIX, pp. 18–18.

[20] LAGUNA, I., AHN, D. H., DE SUPINSKI, B. R., GAMBLIN, T., LEE, G. L., SCHULZ, M., BAGCHI, S., KULKARNI, M., ZHOU, B., CHEN, Z., AND QIN, F. Debugging high-performance computing applications at massive scales. *Communications of the ACM 58*, 9 (2015), 72–81.

[21] LAMPORT, L. Proving the correctness of multiprocess programs. *IEEE Transactions on Software Engineering 3*, 2 (1977), 125–143.

[22] LAMPORT, L. The temporal logic of actions. *ACM Transactions on Programming Languages and Systems 16*, 3 (1994), 872–923.

[23] LAUTERBURG, S., KARMANI, R. K., MARINOV, D., AND AGHA, G. Evaluating ordering heuristics for dynamic partial-order reduction techniques. In *Proceedings of the 13th International Conference on Fundamental Approaches to Software Engineering* (2010), Springer, pp. 308–322.

[24] LEESATAPORNWONGSA, T., HAO, M., JOSHI, P., LUKMAN, J. F., AND GUNAWI, H. S. SAMC: Semantic-aware model checking for fast discovery of deep bugs in cloud systems. In *Proceedings of the 11th USENIX Conference on Operating Systems Design and Implementation* (2014), USENIX, pp. 399–414.

[25] LEIJEN, D., SCHULTE, W., AND BURCKHARDT, S. The design of a task parallel library. In *Proceedings of the 24th ACM SIGPLAN Conference on Object Oriented Programming Systems Languages and Applications* (2009), ACM, pp. 227–242.

[26] MADDOX, P. Testing a distributed system. *ACM Queue 13*, 7 (2015), 10–15.

[27] MUSUVATHI, M., AND QADEER, S. Fair stateless model checking. In *Proceedings of the 29th ACM SIGPLAN Conference on Programming Language Design and Implementation* (2008), ACM, pp. 362–371.

[28] MUSUVATHI, M., QADEER, S., BALL, T., BASLER, G., NAINAR, P. A., AND NEAMTIU, I. Finding and reproducing Heisenbugs in concurrent programs. In *Proceedings of the 8th USENIX Conference on Operating Systems Design and Implementation* (2008), USENIX, pp. 267–280.

[29] NEWCOMBE, C., RATH, T., ZHANG, F., MUNTEANU, B., BROOKER, M., AND DEARDEUFF, M. How amazon web services uses formal methods. *Communications of the ACM 58*, 4 (2015), 66–73.

[30] PNUELI, A. The temporal logic of programs. In *Proceedings of the Foundations of Computer Science* (1977), pp. 46–57.

[31] TASSEY, G. The economic impacts of inadequate infrastructure for software testing. *National Institute of Standards and Technology, Planning Report 02-3* (2002).

[32] THOMSON, P., DONALDSON, A. F., AND BETTS, A. Concurrency testing using schedule bounding: An empirical study. In *Proceedings of the 19th ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming* (2014), ACM, pp. 15–28.

[33] TREYNOR, B. GoogleBlog – Today's outage for several Google services. `http://googleblog.blogspot.com/2014/01/todays-outage-for-several-google.html`, 2014.

[34] ŠIMŠA, J., BRYANT, R., AND GIBSON, G. dBug: Systematic testing of unmodified distributed and multi-threaded systems. In *Proceedings of the 18th International SPIN Conference on Model Checking Software* (2011), Springer, pp. 188–193.

[35] WILCOX, J. R., WOOS, D., PANCHEKHA, P., TATLOCK, Z., WANG, X., ERNST, M. D., AND ANDERSON, T. Verdi: A framework for implementing and formally verifying distributed systems. In *Proceedings of the 36th ACM SIGPLAN Conference on Programming Language Design and Implementation* (2015), ACM, pp. 357–368.

[36] YABANDEH, M., KNEZEVIC, N., KOSTIC, D., AND KUNCAK, V. CrystalBall: Predicting and preventing inconsistencies in deployed distributed systems. In *Proceedings of the 6th USENIX Symposium on Networked Systems Design and Implementation* (2009), USENIX, pp. 229–244.

[37] YANG, J., CHEN, T., WU, M., XU, Z., LIU, X., LIN, H., YANG, M., LONG, F., ZHANG, L., AND ZHOU, L. MODIST: Transparent model checking of unmodified distributed systems. In *Proceedings of the 6th USENIX Symposium on Networked Systems Design and Implementation* (2009), USENIX, pp. 213–228.